

MACHINE LEARNING LAB - TUTORIAL 9

Juan Fernando Espinosa

303158

1. SPAM FILTER USING SVM

A. Build a spam filter using a pre-processed dataset

```
In [1]: from libsvm.python.svmutil import *
import pandas as pd
import pandas as DataFrame
import numpy as np
from nltk import stem
from nltk.corpus import stopwords
from sklearn.feature_extraction import DictVectorizer
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer, TfidfTransformer
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import confusion_matrix
from sklearn.svm import SVC
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
import matplotlib.pyplot as plt
from sklearn.metrics import classification_report
```

```
In [2]: missing_values = ['-','na','Nan','nan','n/a','?']
column_names=['0','1','2','3','4','5','6','7','8','9','10','11','12','13','14','15','16','17','18','19','20','21','22','23','24','25','26','27','28','29','30','31','32','33','34','35','36','37','38','39','40','41','42','43','44','45','46','47','48','49','50','51','52','53','54','55','56','57']
D3 = pd.read_csv("spambase.data", sep=',', na_values = missing_values, names = column_names)
D3.head()
```

Out[2]:

	0	1	2	3	4	5	6	7	8	9	...	48	49	50	51	5
0	0.00	0.64	0.64	0.0	0.32	0.00	0.00	0.00	0.00	0.00	...	0.00	0.000	0.0	0.778	0.00
1	0.21	0.28	0.50	0.0	0.14	0.28	0.21	0.07	0.00	0.94	...	0.00	0.132	0.0	0.372	0.18
2	0.06	0.00	0.71	0.0	1.23	0.19	0.19	0.12	0.64	0.25	...	0.01	0.143	0.0	0.276	0.18
3	0.00	0.00	0.00	0.0	0.63	0.00	0.31	0.63	0.31	0.63	...	0.00	0.137	0.0	0.137	0.00
4	0.00	0.00	0.00	0.0	0.63	0.00	0.31	0.63	0.31	0.63	...	0.00	0.135	0.0	0.135	0.00

5 rows x 58 columns

Training and Testing Splitting

```
In [3]: D3_train = D3.sample(frac=0.7, random_state=1)
D3_test = D3.drop(D3_train.index)

Y_train = D3_train['57'].values
print('y_train set:', len(Y_train))
X_train = D3_train.drop(['57'], axis=1).values
print('X_train set:', len(X_train))
Y_test = D3_test['57'].values
print('y_test set:', len(Y_test))
X_test = D3_test.drop(['57'], axis=1).values
print('X_test set:', len(X_test))
```

```
y_train set: 3221
X_train set: 3221
y_test set: 1380
X_test set: 1380
```

Mixing column name and value into Dictionary - Training

```
In [4]: x = []
for i in range(len(X_train)):
    dictionary = {}
    for j in range(len(X_train[i])):
        dictionary[j]=X_train[i][j]
    x.append(dictionary)
print('First row printing:', '\n', '\n', x[0:1])
```

First row printing:

```
[{0: 0.0, 1: 0.0, 2: 0.0, 3: 0.0, 4: 1.01, 5: 0.0, 6: 0.0, 7: 0.0,
8: 0.0, 9: 0.0, 10: 0.0, 11: 0.0, 12: 0.0, 13: 0.0, 14: 0.0, 15:
0.0, 16: 0.0, 17: 0.0, 18: 0.0, 19: 0.0, 20: 0.0, 21: 0.0, 22: 0.0,
23: 0.0, 24: 0.0, 25: 0.0, 26: 0.0, 27: 0.0, 28: 0.0, 29: 0.0, 30: 0.0,
31: 0.0, 32: 0.0, 33: 0.0, 34: 0.0, 35: 0.0, 36: 5.05, 37: 0.0,
38: 0.0, 39: 0.0, 40: 0.0, 41: 0.0, 42: 0.0, 43: 0.0, 44: 0.0,
45: 0.0, 46: 0.0, 47: 0.0, 48: 0.0, 49: 0.08800000000000001, 50:
0.0, 51: 0.0, 52: 0.08800000000000001, 53: 0.0, 54: 6.718, 55: 33.0,
56: 215.0}]
```

Mixing column name and value into Dictionary - Testing

```
In [5]: x_test = []
        for i in range(len(X_test)):
            dictionary = {}
            for j in range(len(X_test[i])):
                dictionary[j]=X_test[i][j]
            x_test.append(dictionary)
        print('First row printing:', '\n', '\n', x_test[0:1])
```

First row printing:

```
[{0: 0.06, 1: 0.0, 2: 0.71, 3: 0.0, 4: 1.23, 5: 0.19, 6: 0.19, 7:
0.12, 8: 0.64, 9: 0.25, 10: 0.38, 11: 0.45, 12: 0.12, 13: 0.0, 14:
1.75, 15: 0.06, 16: 0.06, 17: 1.03, 18: 1.36, 19: 0.32, 20: 0.51,
21: 0.0, 22: 1.16, 23: 0.06, 24: 0.0, 25: 0.0, 26: 0.0, 27: 0.0, 2
8: 0.0, 29: 0.0, 30: 0.0, 31: 0.0, 32: 0.0, 33: 0.0, 34: 0.0, 35:
0.0, 36: 0.0, 37: 0.0, 38: 0.0, 39: 0.06, 40: 0.0, 41: 0.0, 42: 0.
12, 43: 0.0, 44: 0.06, 45: 0.06, 46: 0.0, 47: 0.0, 48: 0.01, 49: 0
.14300000000000002, 50: 0.0, 51: 0.276, 52: 0.184, 53: 0.01, 54: 9
.821, 55: 485.0, 56: 2259.0}]
```

Training and testing of the model - Classification task

```
In [6]: ranges = np.arange(0.1,100,1)
        len(ranges)
        accuracy_Array = []
        mse_array = []
        prob = svm_problem(Y_train, x, isKernel=True)
        for i in ranges:
            param = svm_parameter('-c '+str(i))
            m = svm_train(prob, param)
            p_label, p_acc, p_val = svm_predict(Y_test, x_test, m)
            ACC, MSE, SCC = evaluations(Y_test, p_label)
            print('MSE = ', MSE)
            accuracy_Array.append(ACC)
            mse_array.append(MSE)
```

```
Accuracy = 75.2899% (1039/1380) (classification)
MSE = 0.24710144927536232
Accuracy = 82.3188% (1136/1380) (classification)
MSE = 0.17681159420289855
Accuracy = 84.4928% (1166/1380) (classification)
MSE = 0.15507246376811595
Accuracy = 84.7101% (1169/1380) (classification)
MSE = 0.15289855072463768
Accuracy = 84.7826% (1170/1380) (classification)
MSE = 0.15217391304347827
Accuracy = 84.9275% (1172/1380) (classification)
MSE = 0.15072463768115943
Accuracy = 84.5652% (1167/1380) (classification)
MSE = 0.15434782608695652
Accuracy = 84.6377% (1168/1380) (classification)
```

```
MSE = 0.1536231884057971
Accuracy = 84.5652% (1167/1380) (classification)
MSE = 0.15434782608695652
Accuracy = 84.6377% (1168/1380) (classification)
MSE = 0.1536231884057971
Accuracy = 84.6377% (1168/1380) (classification)
MSE = 0.1536231884057971
Accuracy = 84.4203% (1165/1380) (classification)
MSE = 0.15579710144927536
Accuracy = 84.4928% (1166/1380) (classification)
MSE = 0.15507246376811595
Accuracy = 84.3478% (1164/1380) (classification)
MSE = 0.1565217391304348
Accuracy = 84.4203% (1165/1380) (classification)
MSE = 0.15579710144927536
Accuracy = 84.3478% (1164/1380) (classification)
MSE = 0.1565217391304348
Accuracy = 84.4203% (1165/1380) (classification)
MSE = 0.15579710144927536
Accuracy = 84.4203% (1165/1380) (classification)
MSE = 0.15579710144927536
Accuracy = 84.4928% (1166/1380) (classification)
MSE = 0.15507246376811595
Accuracy = 84.4928% (1166/1380) (classification)
MSE = 0.15507246376811595
Accuracy = 84.4928% (1166/1380) (classification)
MSE = 0.15507246376811595
Accuracy = 84.4928% (1166/1380) (classification)
MSE = 0.15507246376811595
Accuracy = 84.3478% (1164/1380) (classification)
MSE = 0.1565217391304348
Accuracy = 84.3478% (1164/1380) (classification)
MSE = 0.1565217391304348
Accuracy = 84.3478% (1164/1380) (classification)
MSE = 0.1565217391304348
Accuracy = 84.3478% (1164/1380) (classification)
MSE = 0.1565217391304348
Accuracy = 84.4203% (1165/1380) (classification)
MSE = 0.15579710144927536
Accuracy = 84.4203% (1165/1380) (classification)
MSE = 0.15579710144927536
Accuracy = 84.4203% (1165/1380) (classification)
MSE = 0.15579710144927536
Accuracy = 84.4203% (1165/1380) (classification)
MSE = 0.15579710144927536
Accuracy = 84.3478% (1164/1380) (classification)
MSE = 0.1565217391304348
Accuracy = 84.3478% (1164/1380) (classification)
MSE = 0.1565217391304348
Accuracy = 84.2754% (1163/1380) (classification)
MSE = 0.1572463768115942
Accuracy = 84.2754% (1163/1380) (classification)
MSE = 0.1572463768115942
Accuracy = 84.2754% (1163/1380) (classification)
```

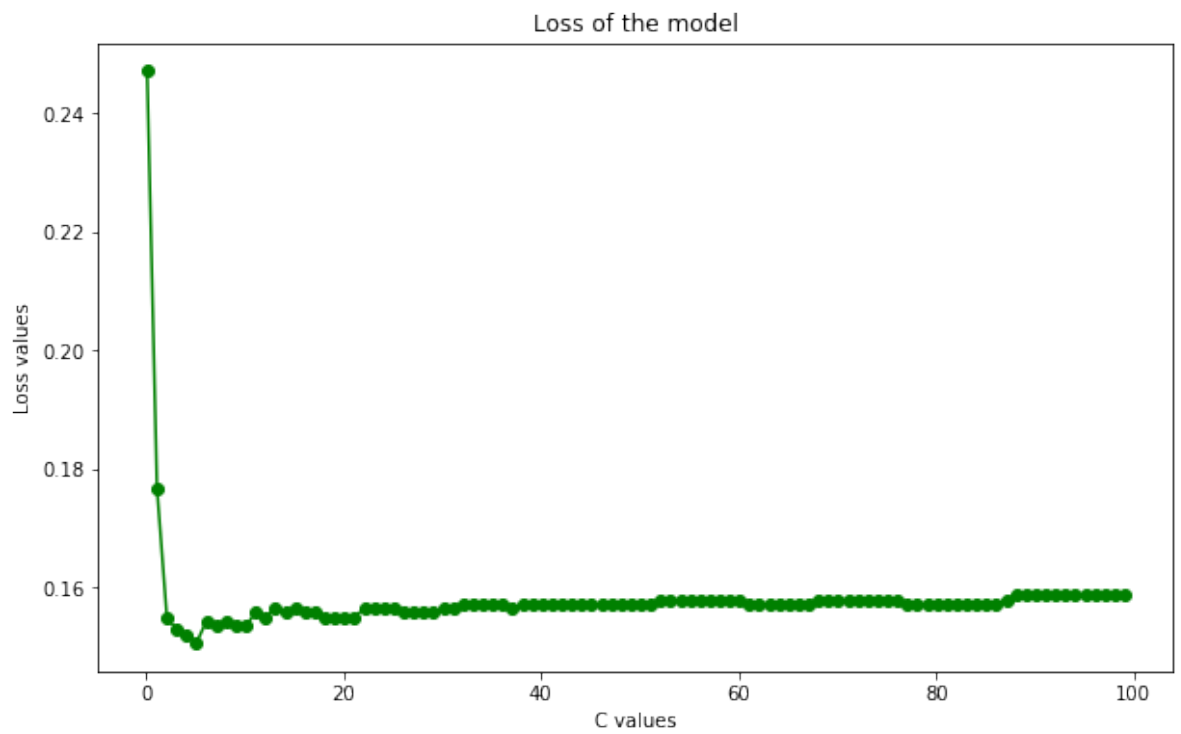
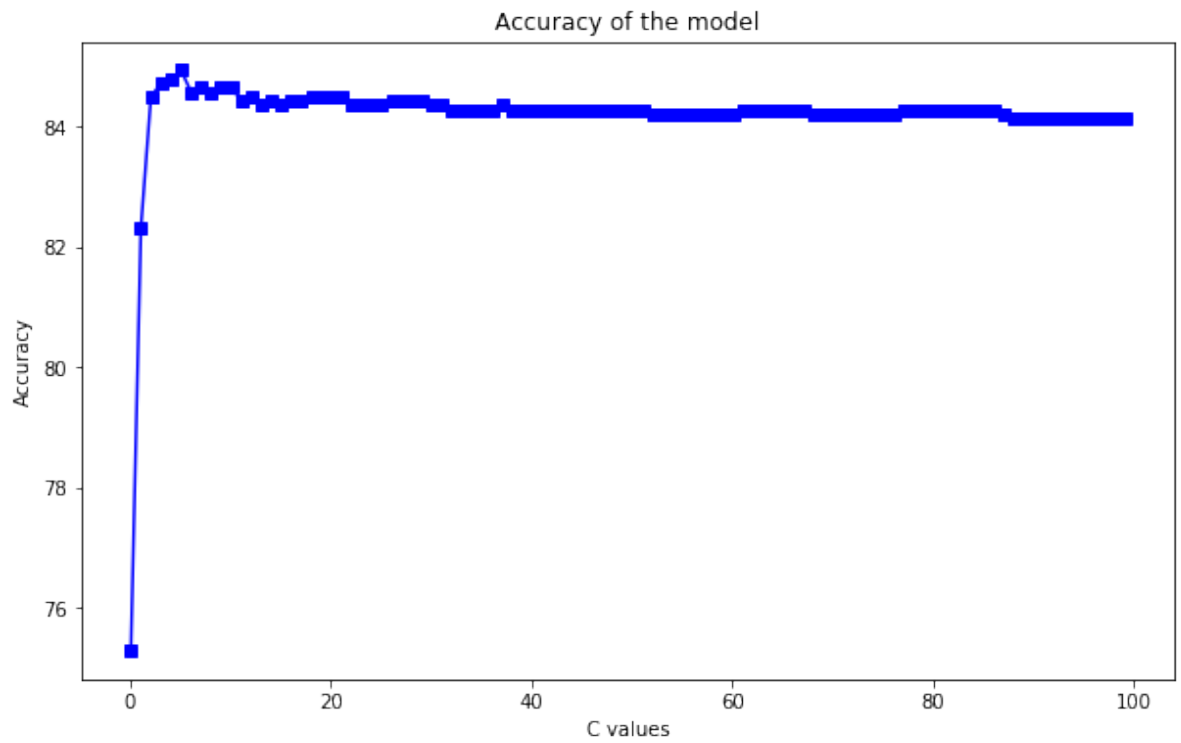
Página 6 de 20

```
MSE = 0.1572463768115942
Accuracy = 84.2754% (1163/1380) (classification)
MSE = 0.1572463768115942
Accuracy = 84.2754% (1163/1380) (classification)
MSE = 0.1572463768115942
Accuracy = 84.2754% (1163/1380) (classification)
MSE = 0.1572463768115942
Accuracy = 84.2754% (1163/1380) (classification)
MSE = 0.1572463768115942
Accuracy = 84.2754% (1163/1380) (classification)
MSE = 0.1572463768115942
Accuracy = 84.2754% (1163/1380) (classification)
MSE = 0.1572463768115942
Accuracy = 84.2029% (1162/1380) (classification)
MSE = 0.15797101449275364
Accuracy = 84.2029% (1162/1380) (classification)
MSE = 0.15797101449275364
Accuracy = 84.2029% (1162/1380) (classification)
MSE = 0.15797101449275364
Accuracy = 84.2029% (1162/1380) (classification)
MSE = 0.15797101449275364
Accuracy = 84.2029% (1162/1380) (classification)
MSE = 0.15797101449275364
Accuracy = 84.2029% (1162/1380) (classification)
MSE = 0.15797101449275364
Accuracy = 84.2029% (1162/1380) (classification)
MSE = 0.15797101449275364
Accuracy = 84.2029% (1162/1380) (classification)
MSE = 0.15797101449275364
Accuracy = 84.2029% (1162/1380) (classification)
MSE = 0.15797101449275364
Accuracy = 84.2754% (1163/1380) (classification)
MSE = 0.1572463768115942
Accuracy = 84.2754% (1163/1380) (classification)
MSE = 0.1572463768115942
Accuracy = 84.2754% (1163/1380) (classification)
MSE = 0.1572463768115942
Accuracy = 84.2754% (1163/1380) (classification)
MSE = 0.1572463768115942
Accuracy = 84.2754% (1163/1380) (classification)
MSE = 0.1572463768115942
Accuracy = 84.2754% (1163/1380) (classification)
MSE = 0.1572463768115942
Accuracy = 84.2754% (1163/1380) (classification)
MSE = 0.1572463768115942
Accuracy = 84.2754% (1163/1380) (classification)
MSE = 0.1572463768115942
Accuracy = 84.2754% (1163/1380) (classification)
MSE = 0.1572463768115942
Accuracy = 84.2754% (1163/1380) (classification)
MSE = 0.1572463768115942
Accuracy = 84.2754% (1163/1380) (classification)
MSE = 0.1572463768115942
Accuracy = 84.2754% (1163/1380) (classification)
MSE = 0.1572463768115942
Accuracy = 84.2754% (1163/1380) (classification)
MSE = 0.1572463768115942
Accuracy = 84.2029% (1162/1380) (classification)
MSE = 0.15797101449275364
Accuracy = 84.1304% (1161/1380) (classification)
```

Plot Accuracy and MSE outputs

Página 8 de 20

Out[7]: <function matplotlib.pyplot.show(*args, **kw)>



```
In [8]: target_names = ['ham', 'spam']
print(classification_report(Y_test, p_label, target_names=target_names))
```

	precision	recall	f1-score	support
ham	0.89	0.84	0.87	850
spam	0.77	0.84	0.80	530
accuracy			0.84	1380
macro avg	0.83	0.84	0.83	1380
weighted avg	0.85	0.84	0.84	1380

B. Pre-processed a dataset and learn SVM

```
In [9]: messages = pd.read_csv('smsspamcollection/SMSSpamCollection', sep='
\t', names=["label", "message"])
print(messages.head())
```

	label	message
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...

```
In [10]: messages.groupby('label').describe()
```

Out[10]:

	message			
	count	unique	top	freq
label				
ham	4825	4516	Sorry, I'll call later	30
spam	747	653	Please call our customer service representativ...	4

```
In [11]: messages['length'] = messages['message'].map(lambda text: len(text))
print(messages.head())
```

	label	message	length
0	ham	Go until jurong point, crazy.. Available only ...	111
1	ham	Ok lar... Joking wif u oni...	29
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	155
3	ham	U dun say so early hor... U c already then say...	49
4	ham	Nah I don't think he goes to usf, he lives aro...	61

```
In [12]: # Lower case for all the words
messages['message'] = messages['message'].map(lambda x: x.lower())
messages['message'].head()
```

```
Out[12]: 0    go until jurong point, crazy.. available only ...
1              ok lar... joking wif u oni...
2    free entry in 2 a wkly comp to win fa cup fina...
3    u dun say so early hor... u c already then say...
4    nah i don't think he goes to usf, he lives aro...
Name: message, dtype: object
```

```
In [13]: # Erasing punctuation of the strings
messages["message"] = messages['message'].str.replace('[^\w\s]', '')
messages["message"].head()
```

```
Out[13]: 0    go until jurong point crazy available only in ...
1              ok lar joking wif u oni
2    free entry in 2 a wkly comp to win fa cup fina...
3          u dun say so early hor u c already then say
4    nah i dont think he goes to usf he lives aroun...
Name: message, dtype: object
```

Streamline workload with pipeline

```
In [14]: pipeline = Pipeline([
            ('vect', CountVectorizer()),
            ('tfidf', TfidfTransformer()),
            ('SVM', SVC())
        ])
```

Slit dataset into train and test

```
In [15]: # Splitting into Train and Test sets.
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(messages['message'],
messages['label'], test_size = 0.3, random_state = 1)
#X_train, X_test, y_train, y_test = train_test_split(messages['message'],
messages['label'], test_size = 0.1, random_state = 1)
print("Original set", messages.shape[0], "observations")
print ("Training set", X_train.shape[0], "observations")
print ("Testing set", X_test.shape[0], "observations")
```

Original set 5572 observations
Training set 3900 observations
Testing set 1672 observations

Parameters to run into the model

```
In [16]: ranges = np.arange(0.1,20,1)
parameters = {'SVM__kernel':['linear', 'rbf'],'SVM__C':[0.0001,0.001,0.1,10,50,100], 'SVM__gamma':[0.1,0.01, 0.001]}

print('SVM__kernel | CV: [linear, rbf]','\n', 'C parameters | CV: [0.0001,0.001,0.1,1,10,50,100]','\n', 'Gama parameters | CV: [0.1,0.01, 0.001]') )
```

SVM__kernel | CV: [linear, rbf]
C parameters | CV: [0.0001,0.001,0.1,1,10,50,100]
Gama parameters | CV: [0.1,0.01, 0.001]

Cross Validation over parameters previously stated

```
In [17]: grid = GridSearchCV(pipeline, param_grid=parameters, cv=5)
```

```
In [18]: # Fit of the trainin
grid.fit(X_train, y_train)
y_hat =grid.fit(X_train, y_train)
print("Best parameter | Training (CV score=%0.3f):" % grid.best_score_)
print(grid.best_params_)
print("Test score = %3.2f" %(grid.score(X_test,y_test)))
```

Best parameter | Training (CV score=0.984):
{'SVM__C': 100, 'SVM__gamma': 0.01, 'SVM__kernel': 'rbf'}
Test score = 0.98

Plot of the relationship between parameters

```

In [19]: def plot_results(index='SVM__C', columns='SVM__gamma'):
    """Select two hyperparameters from which we plot the fluctuations"""
    index = 'param_' + index
    columns = 'param_' + columns

    df = pd.DataFrame(grid.cv_results_)

    other = [c for c in df.columns if c[:6] == 'param_']
    other.remove(index)
    other.remove(columns)
    for col in other:
        df = df[df[col] == grid.best_params_[col[6:]]]

    table_mean = df.pivot_table(index=index, columns=columns,
                                values=['mean_test_score'])
    table_std = df.pivot_table(index=index, columns=columns,
                                values=['std_test_score'])

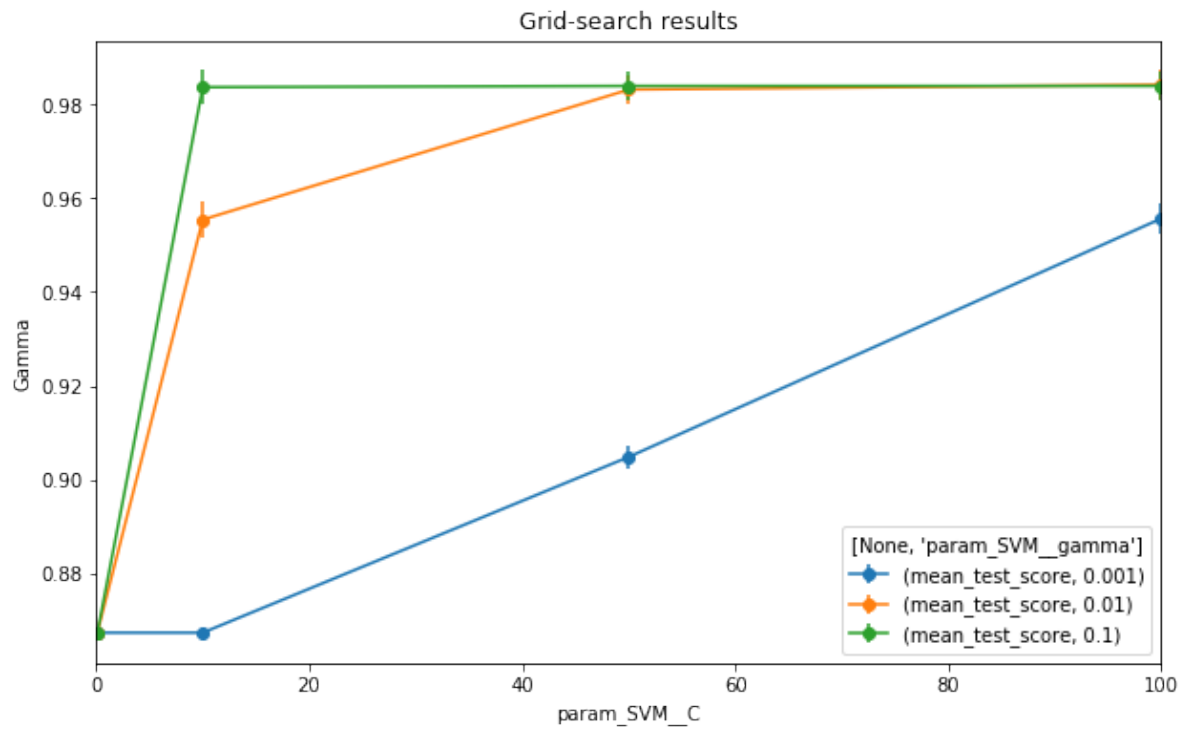
    plt.figure(figsize=(10,6))
    ax = plt.gca()
    for col_mean, col_std in zip(table_mean.columns, table_std.columns):
        table_mean[col_mean].plot(ax=ax, yerr=table_std[col_std], marker='o',
                                label=col_mean)

    plt.title('Grid-search results')
    plt.ylabel('Gamma')
    plt.legend(title=table_mean.columns.names)
    plt.show()

plot_results(index='SVM__C', columns='SVM__gamma')

# SOURCE = []

```



Run the training and test with the optimal values

```
In [20]: parameters = {'SVM__kernel':['rbf'], 'SVM__C':[100], 'SVM__gamma':[0.01]}
grid = GridSearchCV(pipeline, param_grid=parameters, cv=5)
grid.fit(X_train, y_train)

y_true, y_pred = y_test, grid.predict(X_test)
print(classification_report(y_true, y_pred))
p_label, p_acc, p_val = svm_predict(Y_test, x_test, m)
ACC, MSE, SCC = evaluations(Y_test, p_label)
cm = confusion_matrix(y_test, y_pred)
df = pd.DataFrame(cm)
df = df.rename(columns={0: 'Ham', 1: 'Spam'}, index={0: 'Predicted Ham', 1: 'Predicted Spam'})
df
```


	precision	recall	f1-score	support
ham	0.98	1.00	0.99	1442
spam	0.98	0.88	0.93	230
accuracy			0.98	1672
macro avg	0.98	0.94	0.96	1672
weighted avg	0.98	0.98	0.98	1672

Accuracy = 84.1304% (1161/1380) (classification)

Out[20]:

	Ham	Spam
Predicted Ham	1438	4
Predicted Spam	27	203

```
In [21]: plt.matshow(cm)
plt.title('Confusion matrix')
plt.colorbar()
plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.show()
```



2. Compare SVM based spam filter with another model

Streamline workload with pipeline

```
In [22]: logreg=LogisticRegression()
pipeline_lg = Pipeline([
    ('vect',CountVectorizer()),
    ('tfidf',TfidfTransformer()),
    ('clf',logreg)
])
```

Parameters to run into the model

```
In [23]: clf__C=[0.0001,0.001,0.1,10,50,100]
parameter_lg = {'clf__C':[0.0001,0.001,0.1,10,50,100]}
```

Cross Validation over parameters previously stated


```
In [24]: grid_lg = GridSearchCV(pipeline_lg, param_grid=parameter_lg, cv=5)
```

```
In [25]: # Fit of the trainin
grid_lg.fit(X_train, y_train)
print("Best parameter | Training (CV score=%0.3f):" % grid_lg.best_
score_)
print(grid_lg.best_params_)
print("Test score = %3.2f" %(grid_lg.score(X_test,y_test)))

/usr/local/lib/python3.7/site-packages/sklearn/linear_model/logist
ic.py:432: FutureWarning: Default solver will be changed to 'lbfgs
' in 0.22. Specify a solver to silence this warning.
  FutureWarning)

Best parameter | Training (CV score=0.981):
{'clf__C': 100}
Test score = 0.98
```

Run the training and test with the optimal values

```
In [26]: parameters1 = {'clf__C':[100]}
grid_lgl = GridSearchCV(pipeline_lg, param_grid=parameters1, cv=5)
grid_lgl.fit(X_train, y_train)
y_true, y_pred = y_test, grid_lgl.predict(X_test)
print(classification_report(y_true, y_pred))
cm = confusion_matrix(y_test, y_pred)
df = pd.DataFrame(cm)
df = df.rename(columns={0: 'Ham', 1: 'Spam'}, index={0: 'Predicted
Ham', 1: 'Predicted Spam'})
df
```

```
/usr/local/lib/python3.7/site-packages/sklearn/linear_model/logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs'
in 0.22. Specify a solver to silence this warning.
```


```
FutureWarning)
```

	precision	recall	f1-score	support
ham	0.98	1.00	0.99	1442
spam	0.98	0.89	0.93	230
accuracy			0.98	1672
macro avg	0.98	0.94	0.96	1672
weighted avg	0.98	0.98	0.98	1672

Out[26]:

	Ham	Spam
Predicted Ham	1437	5
Predicted Spam	26	204

```
In [27]: plt.matshow(cm,fignum= 1)
plt.title('Confusion matrix')
plt.colorbar()
plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.show()
```



Conclusions

1. The test score does not change in the comparison between the SVM and Logistic models.
2. Regarding to the accuracy, there is no big difference on the outputs considering the confusion matrix and table the SVM has a lightly improvement against the Logistic Model.
3. The dataset structure may be influencing the similarity between the models. Therefore, it is not possible to make conclusions based on the information found.
4. The distribution in train and test dataset influences as well the accuracy for each model.
5. The kernel used in SVM plays a major role: it transforms the data for solving complex problems.

3. Bibliography

<https://www.kaggle.com/darshnaik/sms-spam-or-ham>
(<https://www.kaggle.com/darshnaik/sms-spam-or-ham>)

https://etav.github.io/projects/spam_message_classifier_naive_bayes.html
(https://etav.github.io/projects/spam_message_classifier_naive_bayes.html)

https://radimrehurek.com/data_science_python/
(https://radimrehurek.com/data_science_python/)

<https://www.kaggle.com/darshnaik/sms-spam-or-ham>
(<https://www.kaggle.com/darshnaik/sms-spam-or-ham>)

<https://towardsdatascience.com/a-simple-example-of-pipeline-in-machine-learning-with-scikit-learn-e726ffbb6976> (<https://towardsdatascience.com/a-simple-example-of-pipeline-in-machine-learning-with-scikit-learn-e726ffbb6976>)

https://pactools.github.io/auto_examples/plot_grid_search.html
(https://pactools.github.io/auto_examples/plot_grid_search.html)

<https://blancas.io/sklearn-evaluation/api/plot.html> (<https://blancas.io/sklearn-evaluation/api/plot.html>)